

Check your AWS IAM user access:

aws sts get-caller-identity

```
(base) EKS-autoscaler-demo$ aws sts get-caller-identity
{
  "UserId": "AIDAYGOSJXRYEADY5WBBZ",
  "Account": "5418882672",
  "Arn": "arn:aws:iam::5418882672:user/rajeshv"
}
```

Create EKS Cluster:

```
(base) EKS-autoscaler-demo$ eksctl create cluster --name clusteratoscales --node-type t3.medium --nodes 2 --nodes-min 1 --nodes-max 2 --region us-east-2 --zones=us-east-2a,us-east-2b,us-east-2c
2023-10-29 17:08:02 [i] eksctl version 0.159.0
2023-10-29 17:08:02 [i] using region us-east-2
2023-10-29 17:08:03 [i] subnets for us-east-2a - public:192.168.0.0/19 private:192.168.96.0/19
2023-10-29 17:08:03 [i] subnets for us-east-2b - public:192.168.32.0/19 private:192.168.128.0/19
2023-10-29 17:08:03 [i] subnets for us-east-2c - public:192.168.64.0/19 private:192.168.160.0/19
2023-10-29 17:08:03 [i] nodegroup "ng-d800e9d1" will use "" [AmazonLinux2/1.25]
2023-10-29 17:08:03 [i] using Kubernetes version 1.25
2023-10-29 17:08:03 [i] creating EKS cluster "clusteratoscales" in "us-east-2" region with managed nodes
2023-10-29 17:08:03 [i] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
2023-10-29 17:08:03 [i] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=us-east-2 --cluster=clusteratoscales'
2023-10-29 17:08:03 [i] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "clusteratoscales" in "us-east-2"
2023-10-29 17:08:03 [i] CloudWatch logging will not be enabled for cluster "clusteratoscales" in "us-east-2"
2023-10-29 17:08:03 [i] you can enable it with 'eksctl utils update-cluster-logging --enable-types={SPECIFY-YOUR-LOG-TYPES-HERE (e.g. all)} --region=us-east-2 --cluster=clusteratoscales'
2023-10-29 17:08:03 [i]
2 sequential tasks: { create cluster control plane "clusteratoscales",
  2 sequential sub-tasks: {
    wait for control plane to become ready,
    create managed nodegroup "ng-d800e9d1",
  }
}
2023-10-29 17:08:03 [i] building cluster stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:08:05 [i] deploying stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:08:35 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:09:07 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:10:08 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:11:10 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:12:11 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:13:12 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
2023-10-29 17:14:14 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"
```

Logs:

```
(base) priti@rv:/DATA/Virtual-Python-ENV/EKS-autoscaler-demo$ eksctl create cluster --name clusteratoscales --node-type t3.medium --nodes 2 --nodes-min 1 --nodes-max 2 --region us-east-2 --zones=us-east-2a,us-east-2b,us-east-2c
2023-10-29 17:08:02 [i] eksctl version 0.159.0
2023-10-29 17:08:02 [i] using region us-east-2
2023-10-29 17:08:03 [i] subnets for us-east-2a - public:192.168.0.0/19 private:192.168.96.0/19
2023-10-29 17:08:03 [i] subnets for us-east-2b - public:192.168.32.0/19 private:192.168.128.0/19
2023-10-29 17:08:03 [i] subnets for us-east-2c - public:192.168.64.0/19 private:192.168.160.0/19
2023-10-29 17:08:03 [i] nodegroup "ng-d800e9d1" will use "" [AmazonLinux2/1.25]
2023-10-29 17:08:03 [i] using Kubernetes version 1.25
2023-10-29 17:08:03 [i] creating EKS cluster "clusteratoscales" in "us-east-2" region with managed nodes
2023-10-29 17:08:03 [i] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
```

2023-10-29 17:08:03 [i] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=us-east-2 --cluster=clusteratoscales'

2023-10-29 17:08:03 [i] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "clusteratoscales" in "us-east-2"

2023-10-29 17:08:03 [i] CloudWatch logging will not be enabled for cluster "clusteratoscales" in "us-east-2"

2023-10-29 17:08:03 [i] you can enable it with 'eksctl utils update-cluster-logging --enable-types={SPECIFY-YOUR-LOG-TYPES-HERE (e.g. all)} --region=us-east-2 --cluster=clusteratoscales'

2023-10-29 17:08:03 [i]

```

2 sequential tasks: { create cluster control plane "clusteratoscales",
  2 sequential sub-tasks: {
    wait for control plane to become ready,
    create managed nodegroup "ng-d800e9d1",
  }
}

```

2023-10-29 17:08:03 [i] building cluster stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:08:05 [i] deploying stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:08:35 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:09:07 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:10:08 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:11:10 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:12:11 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:13:12 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:14:14 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:15:15 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:16:16 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:17:18 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-cluster"

2023-10-29 17:19:28 [i] building managed nodegroup stack "eksctl-clusteratoscales-nodegroup-ng-d800e9d1"

2023-10-29 17:19:30 [i] deploying stack "eksctl-clusteratoscales-nodegroup-ng-d800e9d1"

2023-10-29 17:19:30 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-nodegroup-ng-d800e9d1"

2023-10-29 17:20:01 [i] waiting for CloudFormation stack "eksctl-clusteratoscales-nodegroup-ng-d800e9d1"

2023-10-29 17:20:42 [i] waiting for CloudFormation stack
"eksctl-clusteratoscales-nodegroup-ng-d800e9d1"

2023-10-29 17:21:59 [i] waiting for CloudFormation stack
"eksctl-clusteratoscales-nodegroup-ng-d800e9d1"
2023-10-29 17:21:59 [i] waiting for the control plane to become ready
2023-10-29 17:22:01 [✓] saved kubeconfig as "/home/priti/.kube/config"
2023-10-29 17:22:01 [i] no tasks
2023-10-29 17:22:01 [✓] all EKS cluster resources for "clusteratoscales" have been created
2023-10-29 17:22:02 [i] nodegroup "ng-d800e9d1" has 2 node(s)
2023-10-29 17:22:02 [i] node "ip-192-168-59-27.us-east-2.compute.internal" is ready
2023-10-29 17:22:02 [i] node "ip-192-168-70-219.us-east-2.compute.internal" is ready
2023-10-29 17:22:02 [i] waiting for at least 1 node(s) to become ready in "ng-d800e9d1"
2023-10-29 17:22:02 [i] nodegroup "ng-d800e9d1" has 2 node(s)
2023-10-29 17:22:02 [i] node "ip-192-168-59-27.us-east-2.compute.internal" is ready
2023-10-29 17:22:02 [i] node "ip-192-168-70-219.us-east-2.compute.internal" is ready
2023-10-29 17:22:04 [i] kubectl command should work with "/home/priti/.kube/config", try
'kubectl get nodes'
2023-10-29 17:22:04 [✓] EKS cluster "clusteratoscales" in "us-east-2" region is ready

aws

Services

Search

[Alt+S]

Disk size

Select the size of the attached EBS volume for each node.

20

GiB

Node group scaling configuration

Desired size

Set the desired number of nodes that the group should launch with initially.

2

nodes

Desired node size must be greater than or equal to 0

Minimum size

Set the minimum number of nodes that the group can scale in to.

2

nodes

Minimum node size must be greater than or equal to 0

Maximum size

Set the maximum number of nodes that the group can scale out to.

5

nodes

Maximum node size must be greater than or equal to 1 and cannot be lower than the minimum size

Node group update configuration

Info

Maximum unavailable

Set the maximum number or percentage of unavailable nodes to be tolerated during the node group version update.

☒ Number

Enter a number

☐ Percentage

Specify a percentage

Value

1

node

Node count must be greater than 0.

Cancel

Previous

Next

aws Services Search [Alt+S]

EKS > Clusters > clusterautoscales > Add node group

Step 1
[Configure node group](#)

Step 2
[Set compute and scaling configuration](#)

Step 3
Specify networking

Step 4
[Review and create](#)

Specify networking

Node group network configuration

These properties cannot be changed after the node group is created.

Subnets [Info](#)

Specify the subnets in your VPC where your nodes will run. To create a new subnet, go to the corresponding page in the [VPC console](#).

Select subnets

- subnet-0fbc4dc85e534c89e | eksctl-clusterautoscales-cluster/SubnetPublicUSEAST2B
- subnet-0c18023baf81efbd1 | eksctl-clusterautoscales-cluster/SubnetPrivateUSEAST2A
- subnet-0b1d4739e1317f6ac | eksctl-clusterautoscales-cluster/SubnetPublicUSEAST2C
- subnet-04c07f5e264ec9a91 | eksctl-clusterautoscales-cluster/SubnetPrivateUSEAST2B
- subnet-01b74f4b52c4d857f | eksctl-clusterautoscales-cluster/SubnetPrivateUSEAST2C
- subnet-0b4e66049f3e345a8 | eksctl-clusterautoscales-cluster/SubnetPublicUSEAST2A

☐ Configure remote access to nodes [Info](#)

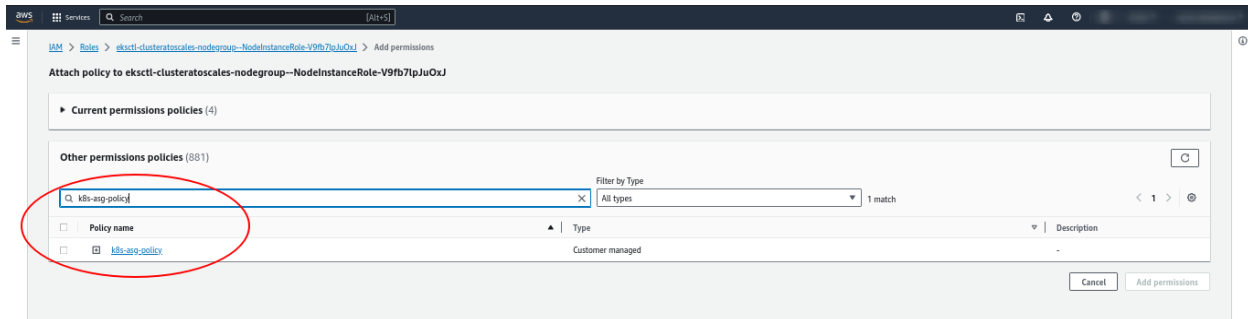
Cancel Previous **Next**

Create IAM policy:

aws iam create-policy --policy-name k8s-asg-policy --policy-document file://k8s-asg-policy.json

```
(base) /EKS-autoscaler-demo$ aws iam create-policy --policy-name k8s-asg-policy --policy-document file://k8s-asg-policy.json
{
  "Policy": {
    "PolicyName": "k8s-asg-policy",
    "PolicyId": "ANPAYGOSJXRYEBMWFJ8B2",
    "Arn": "arn:aws:iam::123456789012:policy/k8s-asg-policy",
    "Path": "/",
    "DefaultVersionId": "v1",
    "AttachmentCount": 0,
    "PermissionsBoundaryUsageCount": 0,
    "IsAttachable": true,
    "CreateDate": "2023-10-29T12:18:27+00:00",
    "UpdateDate": "2023-10-29T12:18:27+00:00"
  }
}
```

Now Go to Cluster >> Configuration >> Node group >> From Detail tab
Click on Role ARN



Deploy the Cluster Autoscaler on Management control plan

Open file and change your cluster name: `vim cluster-autoscaler-autodiscover.yaml`

```
161 - --skip-nodes-with-local-storage=false
162 - --expander=least-waste
163 - --node-group-auto-discovery=asg:tag=k8s.io/cluster-autoscaler/enabled,k8s.io/cluster-autoscaler/clusterautoscales
164 - --balance-similar-node-groups
165 - --skip-nodes-with-system-pods=false
```

Change your cluster name →

Apply the changes:

```
(base) /EKS-autoscaler-demo$ kubectl apply -f cluster-autoscaler-autodiscover.yaml
serviceaccount/cluster-autoscaler created
clusterrole.rbac.authorization.k8s.io/cluster-autoscaler created
role.rbac.authorization.k8s.io/cluster-autoscaler created
clusterrolebinding.rbac.authorization.k8s.io/cluster-autoscaler created
rolebinding.rbac.authorization.k8s.io/cluster-autoscaler created
deployment.apps/cluster-autoscaler created
(base) /EKS-autoscaler-demo$
```

Now we have 4 Nodes with cluster autoscaler pod deployed:

```
(base) /EKS-autoscaler-demo$ kubectl get pods -A
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  aws-node-8rzj4                         1/1     Running   0           37m
kube-system  aws-node-m1rzs                         1/1     Running   0           51m
kube-system  aws-node-swcg9                         1/1     Running   0           51m
kube-system  aws-node-xzpr6                         1/1     Running   0           37m
kube-system  cluster-autoscaler-78c585d7dc-rfvrn    0/1     OOMKilled 2 (54s ago) 2m3s
kube-system  coredns-8fd4db68f-jc6pb               1/1     Running   0           59m
kube-system  coredns-8fd4db68f-pgsgb               1/1     Running   0           59m
kube-system  kube-proxy-544xd                       1/1     Running   0           51m
kube-system  kube-proxy-knd28                       1/1     Running   0           37m
kube-system  kube-proxy-tj1lv                       1/1     Running   0           51m
kube-system  kube-proxy-v72mw                       1/1     Running   0           37m
(base) /EKS-autoscaler-demo$ kubectl get nodes
NAME                                                    STATUS   ROLES    AGE   VERSION
ip-192-168-16-27.us-east-2.compute.internal           Ready    <none>   38m   v1.25.13-eks-43840fb
ip-192-168-187-95.us-east-2.compute.internal           Ready    <none>   38m   v1.25.13-eks-43840fb
ip-192-168-59-27.us-east-2.compute.internal            Ready    <none>   52m   v1.25.13-eks-43840fb
ip-192-168-70-219.us-east-2.compute.internal           Ready    <none>   52m   v1.25.13-eks-43840fb
(base) /EKS-autoscaler-demo$
```

Note: To perform an auto scaling demo, mimicked as the actual traffic on EKS cluster in production environment I created nginx deployment with single pod and scaled it by 100 pods, so then cluster auto scaled out and when traffic decreased EKS cluster scaled in with default nodes.

Created nginx deployment with one pod

```
^C(base) [REDACTED] /EKS-autoscaler-demo$ kubectl create deployment autoscale-deployment-demo --image=nginx
deployment.apps/autoscale-deployment-demo created

(base) [REDACTED] /EKS-autoscaler-demo$ kubectl get deploy
NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
autoscale-deployment-demo          1/1      1              1            15s
(base) [REDACTED] /EKS-autoscaler-demo$
```

Scaled out the deployment with 100 Pods, and watched EKS cluster auto scaling works.


```

autoscale-deployment-demo-9bfc7d8cd-49kgq 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-54b6v 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-577w9 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-5prjp 0/1 Pending 0 29s
autoscale-deployment-demo-9bfc7d8cd-5rlph 1/1 Running 0 30s
autoscale-deployment-demo-9bfc7d8cd-5sjwp 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-5vkdw 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-64gxc 1/1 Running 0 30s
autoscale-deployment-demo-9bfc7d8cd-6jd8m 0/1 ContainerCreating 0 30s
autoscale-deployment-demo-9bfc7d8cd-6stvx 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-7cqgd 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-7f4wn 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-7n5hm 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-7nwjf 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-8qxwz 0/1 ContainerCreating 0 29s
autoscale-deployment-demo-9bfc7d8cd-8rfwb 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-8wt6x 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-bt5n4 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-bxvfs 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-c67t4 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-c7r67 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-cf45v 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-cv6kh 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-cwrp4 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-djw52 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-dm8q9 1/1 Running 0 30s
autoscale-deployment-demo-9bfc7d8cd-dxmwt 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-f7nln 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-fhtls 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-flplc 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-fpz26 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-g8z7h 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-gmzt6 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-gv7j9 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-gxbhn 1/1 Running 0 30s
autoscale-deployment-demo-9bfc7d8cd-h5hsb 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-h96mg 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-hqwbm 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-hrfpp 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-htx9s 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-jjfn2 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-jtkzv 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-kdgj6 0/1 Pending 0 28s
autoscale-deployment-demo-9bfc7d8cd-kg8ml 0/1 Pending 0 26s
autoscale-deployment-demo-9bfc7d8cd-kp29t 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-ksg8r 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-l4mdp 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-l5887 1/1 Running 0 29s
autoscale-deployment-demo-9bfc7d8cd-lznj6 0/1 Pending 0 27s
autoscale-deployment-demo-9bfc7d8cd-ddwb 0/1 Pending 0 27s

```

When traffic increases EKS cluster scaled up:

```

(base) EKS-autoscaler-demo$ kubectl get nodes -w
NAME                                STATUS    ROLES    AGE    VERSION
ip-192-168-16-27.us-east-2.compute.internal Ready    <none>   47m    v1.25.13-eks-43840fb
ip-192-168-187-95.us-east-2.compute.internal Ready    <none>   47m    v1.25.13-eks-43840fb
ip-192-168-59-27.us-east-2.compute.internal Ready    <none>   61m    v1.25.13-eks-43840fb
ip-192-168-70-219.us-east-2.compute.internal Ready    <none>   61m    v1.25.13-eks-43840fb
ip-192-168-16-27.us-east-2.compute.internal Ready    <none>   48m    v1.25.13-eks-43840fb
ip-192-168-187-95.us-east-2.compute.internal Ready    <none>   48m    v1.25.13-eks-43840fb
ip-192-168-59-27.us-east-2.compute.internal Ready    <none>   62m    v1.25.13-eks-43840fb
^C(base) /EKS-autoscaler-demo$ date
Sunday 29 October 2023 06:24:21 PM IST
(base) /EKS-autoscaler-demo$

```

When traffic decreased EKS cluster scaled down:

```
Sunday 29 October 2023 06:24:21 PM IST
(base) /EKS-autoscaler-demo$ kubectl get nodes -w
NAME                                STATUS    ROLES    AGE   VERSION
ip-192-168-16-27.us-east-2.compute.internal Ready    <none>   53m   v1.25.13-eks-43840fb
ip-192-168-187-95.us-east-2.compute.internal Ready    <none>   53m   v1.25.13-eks-43840fb
ip-192-168-59-27.us-east-2.compute.internal Ready    <none>   67m   v1.25.13-eks-43840fb
ip-192-168-70-219.us-east-2.compute.internal Ready    <none>   67m   v1.25.13-eks-43840fb
```

After demo:

##eksctl delete cluster --name clusterat scales --region us-east-2